

Lecture12 Generative Models

sgc

March 24, 2022

1 Supervised vs Unsupervised Learning

Supervised Learning : Sample (data + labels) Examples: Classification, regression, object detection, semantic segmentation, image captioning, etc.

Unsupervised Learning : Sample (data, no labels) Learn Hidden Structure
Examples: Clustering, dimensionality reduction, feature learning, density estimation, etc.

2 Generative Modeling

Given training data. generate new samples from same distribution: Learn p_{model} that approximates p_{data} .

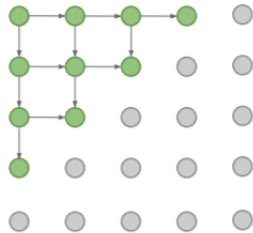
Can learn important features to reconstruct data set.

3 PixelRNN and PixelCNN

FVBN: maximize likelihood of training data : $p(x) = \prod_{i=1}^n p(x_i | x_1, \dots, x_{i-1})$ x_i is pixel No.i.

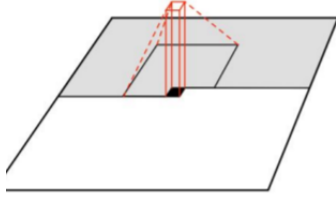
Explicit Model: Explicitly compute likelihood.

PixelRNN generate image start from the corner. Depend on previous pixels.



But it is slow.

PixelCNN: Still start from corner, but now using a CNN over context region.




Relatively faster but still slow

4 Variational Autoencoders VAE

Autoencoders reconstruct data through latent z that represents the most meaningful features.

Autoencoders can be used to make a supervised model more efficient.

Variational Autoencoders:



x is an image, and z is latent factors to generate x : Choose $p(z)$ to be simple, $p(x|z)$ is complex (neural network).

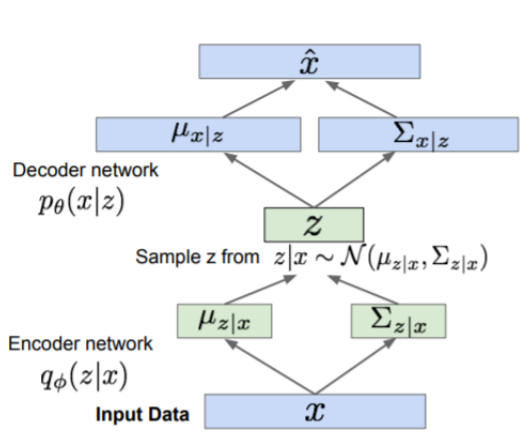
Maximize likelihood: $p_\theta(x) = \int p_\theta(z)p_\theta(x|z)$, but it is intractable.

Use distribution $q_\phi(x|z)$ to approximate $p_\theta(x|z)$

$$\begin{aligned}
 \log p_\theta(x^{(i)}) &= \mathbf{E}_{z \sim q_\phi(z|x^{(i)})} [\log p_\theta(x^{(i)})] && (p_\theta(x^{(i)} \text{ Does not depend on } z) \\
 &= \mathbf{E}_z \left[\log \frac{p_\theta(x^{(i)} | z) p_\theta(z)}{p_\theta(z | x^{(i)})} \right] && (\text{Bayes' Rule}) \\
 \text{Decoder:} &= \mathbf{E}_z \left[\log \frac{p_\theta(x^{(i)} | z) p_\theta(z) q_\phi(z | x^{(i)})}{p_\theta(z | x^{(i)}) q_\phi(z | x^{(i)})} \right] && (\text{Multiply by constant}) \\
 \text{reconstruct} &= \mathbf{E}_z \left[\log p_\theta(x^{(i)} | z) \right] - \mathbf{E}_z \left[\log \frac{q_\phi(z | x^{(i)})}{p_\theta(z)} \right] + \mathbf{E}_z \left[\log \frac{q_\phi(z | x^{(i)})}{p_\theta(z | x^{(i)})} \right] && (\text{Logarithms}) \\
 \text{the input data} &= \underbrace{\mathbf{E}_z \left[\log p_\theta(x^{(i)} | z) \right] - D_{KL}(q_\phi(z | x^{(i)}) || p_\theta(z))}_{\mathcal{L}(x^{(i)}, \theta, \phi)} + \underbrace{D_{KL}(q_\phi(z | x^{(i)}) || p_\theta(z | x^{(i)}))}_{\geq 0} \\
 &\text{Tractable lower bound which we can take} \\
 &\text{gradient of and optimize! } (p_\theta(x|z) \text{ differentiable,} \\
 &\text{KL term differentiable)}
 \end{aligned}$$

Encoder: make approximate posterior distribution close to prior

Optimizing TLB.



Encoder network is used to lessen KL divergence, and Decoder network is used to maximize Expectation. For every minibatch of input data: compute this forward pass, and then backprop.

After training, just sample z from prior and use decoder network.

5 Generative Adversarial Networks (GANs)

GANs is a Two-player game consists of the Discriminator network that tries to distinguish between real and fake images and the Generator network that tries to fool the discriminator by generating real-looking images.

Objective function:

Discriminator outputs likelihood in (0,1) of real image

Minimax objective function:

$$\min_{\theta_g} \max_{\theta_d} \left[\mathbb{E}_{x \sim p_{data}} \log \underbrace{D_{\theta_d}(x)}_{\substack{\text{Discriminator output} \\ \text{for real data } x}} + \mathbb{E}_{z \sim p(z)} \log(1 - \underbrace{D_{\theta_d}(G_{\theta_g}(z))}_{\substack{\text{Discriminator output for} \\ \text{generated fake data } G(z)}}) \right]$$

- Discriminator (θ_d) wants to **maximize objective** such that $D(x)$ is close to 1 (real) and $D(G(z))$ is close to 0 (fake)
- Generator (θ_g) wants to **minimize objective** such that $D(G(z))$ is close to 1 (discriminator is fooled into thinking generated $G(z)$ is real)

Alternate between:

1. **Gradient ascent** on discriminator

$$\max_{\theta_d} \left[\mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

2. **Instead: Gradient ascent** on generator, **different objective**

$$\max_{\theta_g} \mathbb{E}_{z \sim p(z)} \log(D_{\theta_d}(G_{\theta_g}(z)))$$



Use Gradient ascent get higher gradient for bad samples and works better.

```

for number of training iterations do
  for  $k$  steps do
    • Sample minibatch of  $m$  noise samples  $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$  from noise prior  $p_g(\mathbf{z})$ .
    • Sample minibatch of  $m$  examples  $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$  from data generating distribution  $p_{\text{data}}(\mathbf{x})$ .
    • Update the discriminator by ascending its stochastic gradient:
      
$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D_{\theta_d}(\mathbf{x}^{(i)}) + \log(1 - D_{\theta_d}(G_{\theta_g}(\mathbf{z}^{(i)}))) \right]$$

    end for
    • Sample minibatch of  $m$  noise samples  $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$  from noise prior  $p_g(\mathbf{z})$ .
    • Update the generator by ascending its stochastic gradient (improved objective):
      
$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log(D_{\theta_d}(G_{\theta_g}(\mathbf{z}^{(i)})))$$

  end for

```

Some find $k=1$ more stable.