

# Lecture12 Generative Models

sgc

March 27, 2022

## 1 Supervised vs Unsupervised Learning

Supervised Learning : Sample (data + labels) Examples: Classification, regression, object detection, semantic segmentation, image captioning, etc.

Unsupervised Learning : Sample (data, no labels ) Learn Hidden Structure  
Examples: Clustering, dimensionality reduction, feature learning, density estimation, etc.

## 2 Generative Modeling

Given training data. generate new samples from same distribution: Learn  $p_{model}$  that approximates  $p_{data}$ .

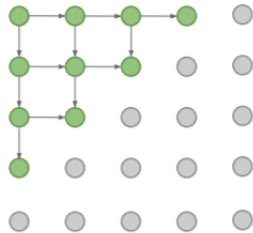
Can learn important features to reconstruct data set.

## 3 PixelRNN and PixelCNN

FVBN: maximize likelihood of training data :  $p(x) = \prod_{i=1}^n p(x_i | x_1, \dots, x_{i-1})$   $x_i$  is pixel No.i.

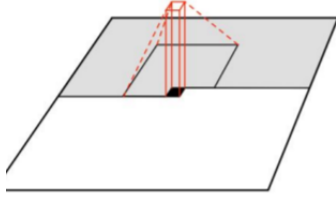
Explicit Model: Explicitly compute likelihood.

PixelRNN generate image start from the corner. Depend on previous pixels.



But it is slow.

PixelCNN: Still start from corner, but now using a CNN over context region.



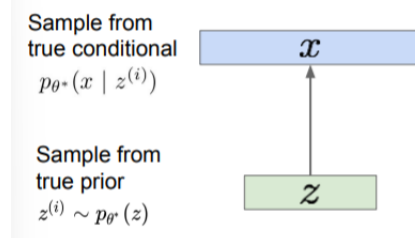
Relatively faster but still slow

## 4 Variational Autoencoders VAE

Autoencoders reconstruct data through latent  $z$  that represents the most meaningful features.

Autoencoders can be used to make a supervised model more efficient.

Variational Autoencoders:



$x$  is an image, and  $z$  is latent factors to generate  $x$ : Choose  $p(z)$  to be simple,  $p(x|z)$  is complex (neural network).

Maximize likelihood:  $p_{\theta}(x) = \int p_{\theta}(z)p_{\theta}(x|z)$ , but it is intractable.

Use distribution  $q_{\phi}(x|z)$  to approximate  $p_{\theta}(x|z)$

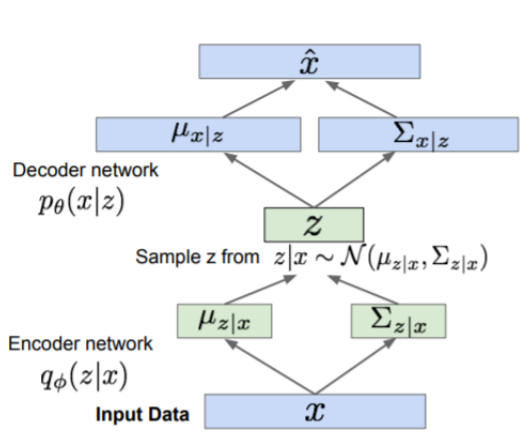
$$\begin{aligned}
 \log p_{\theta}(x^{(i)}) &= \mathbf{E}_{z \sim q_{\phi}(z|x^{(i)})} [\log p_{\theta}(x^{(i)})] \quad (p_{\theta}(x^{(i)}) \text{ Does not depend on } z) \\
 &= \mathbf{E}_z \left[ \log \frac{p_{\theta}(x^{(i)} | z) p_{\theta}(z)}{p_{\theta}(z | x^{(i)})} \right] \quad (\text{Bayes' Rule}) \\
 &= \mathbf{E}_z \left[ \log \frac{p_{\theta}(x^{(i)} | z) p_{\theta}(z) q_{\phi}(z | x^{(i)})}{p_{\theta}(z | x^{(i)}) q_{\phi}(z | x^{(i)})} \right] \quad (\text{Multiply by constant}) \\
 &= \mathbf{E}_z \left[ \log p_{\theta}(x^{(i)} | z) \right] - \mathbf{E}_z \left[ \log \frac{q_{\phi}(z | x^{(i)})}{p_{\theta}(z)} \right] + \mathbf{E}_z \left[ \log \frac{q_{\phi}(z | x^{(i)})}{p_{\theta}(z | x^{(i)})} \right] \quad (\text{Logarithms}) \\
 &= \underbrace{\mathbf{E}_z \left[ \log p_{\theta}(x^{(i)} | z) \right] - D_{KL}(q_{\phi}(z | x^{(i)}) || p_{\theta}(z))}_{\mathcal{L}(x^{(i)}, \theta, \phi)} + \underbrace{D_{KL}(q_{\phi}(z | x^{(i)}) || p_{\theta}(z | x^{(i)}))}_{\geq 0}
 \end{aligned}$$

Decoder: reconstruct the input data

Encoder: make approximate posterior distribution close to prior

Tractable lower bound which we can take gradient of and optimize! ( $p_{\theta}(x|z)$  differentiable, KL term differentiable)

Optimizing ELBO.



Encoder network is used to lessen KL divergence, and Decoder network is used to maximize Expectation. For every minibatch of input data: compute this forward pass, and then backprop.

After training, just sample  $z$  from prior and use decoder network.

## 5 Generative Adversarial Networks (GANs)

GANs is a Two-player game consists of the Discriminator network that tries to distinguish between real and fake images and the Generator network that tries to fool the discriminator by generating real-looking images.

Objective function:

Discriminator outputs likelihood in (0,1) of real image

Minimax objective function:

$$\min_{\theta_g} \max_{\theta_d} \left[ \mathbb{E}_{x \sim p_{data}} \log \underbrace{D_{\theta_d}(x)}_{\substack{\text{Discriminator output} \\ \text{for real data } x}} + \mathbb{E}_{z \sim p(z)} \log(1 - \underbrace{D_{\theta_d}(G_{\theta_g}(z))}_{\substack{\text{Discriminator output for} \\ \text{generated fake data } G(z)}}) \right]$$

- Discriminator ( $\theta_d$ ) wants to **maximize objective** such that  $D(x)$  is close to 1 (real) and  $D(G(z))$  is close to 0 (fake)
- Generator ( $\theta_g$ ) wants to **minimize objective** such that  $D(G(z))$  is close to 1 (discriminator is fooled into thinking generated  $G(z)$  is real)

Alternate between:

1. **Gradient ascent** on discriminator

$$\max_{\theta_d} \left[ \mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

2. **Instead: Gradient ascent** on generator, **different objective**

$$\max_{\theta_g} \mathbb{E}_{z \sim p(z)} \log(D_{\theta_d}(G_{\theta_g}(z)))$$



Use Gradient ascent get higher gradient for bad samples and works better.

```

for number of training iterations do
  for  $k$  steps do
    • Sample minibatch of  $m$  noise samples  $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$  from noise prior  $p_g(\mathbf{z})$ .
    • Sample minibatch of  $m$  examples  $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$  from data generating distribution  $p_{\text{data}}(\mathbf{x})$ .
    • Update the discriminator by ascending its stochastic gradient:
      
$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D_{\theta_d}(x^{(i)}) + \log(1 - D_{\theta_d}(G_{\theta_g}(z^{(i)}))) \right]$$

    end for
    • Sample minibatch of  $m$  noise samples  $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$  from noise prior  $p_g(\mathbf{z})$ .
    • Update the generator by ascending its stochastic gradient (improved objective):
      
$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log(D_{\theta_d}(G_{\theta_g}(z^{(i)})))$$

  end for

```

Some find  $k=1$  more stable.